

# Terrain Modelling with Generative AI

Authors: Ewan Walker, Harry Fletcher, Ed Warneken

Supervisor: Dr Nick Keeper

MAS3091 2023-2024 – Group 8



## Overview

A detailed outline of the objectives are as follows:

1. Code multiple of these methods for analysis by varying parameters and create plots, colourmaps and 3D projections (where applicable). Implement reproducibility via the use of seeds (an initialisation of pseudorandom number generators to recreate the same 'random' terrain). Utilise noise in heightmaps – fluently blended 'random' normalised values (ranging from 0 to 1) sorted in matrices to create 3D projections using the values as height components.
2. Explore principal component analysis to identify the most important data-points representing the terrain and eliminate any excess for dimensionality reduction, thereby compressing the volume of data. Find out the optimum number of principle components to minimize loss of quality when compared to the original terrain.
3. Build autoencoders, trained on the generated heightmaps, to create a 2D grid in which clusters of points correlate to specific features of the terrain. These are much more efficient than traditional methods.
4. Analyse real-world LiDAR scans via the autoencoders to create realistic, new, predictable terrain from a few select coordinates.

Clearly, that is a lot to unpack! However, the simplistic beauty of the process can be quickly illustrated, and all the fancy concepts and definitions are reduced to an easily understandable format.

Firstly, to understand what we mean by 'terrain', take a widely recognised example: Minecraft. In Minecraft, vast, unique worlds are near-instantaneously spawned into existence for the player to explore. They are random; they also contain consistently predictable patterns. One can reproduce infinitely generating, complex random worlds using a short string of numbers, a seed, whilst no two seeds are the same. How can this be? Below is information to lay a bedrock understanding of all these processes.


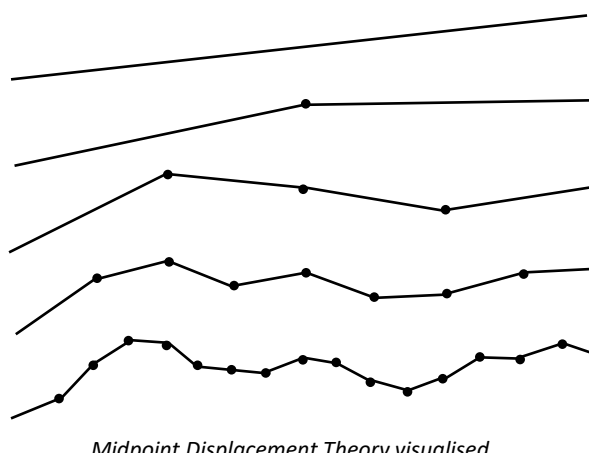


[1]Minecraft Landscape: a prime example of terrain generation

## Standard Methods

### Midpoint Displacement


The Midpoint Displacement method begins with taking two end points of a 1D line and iteratively calculates the midpoint of the connecting line, displacing each point in the y-direction with an incrementally reduced range of randomness (scaled by a h parameter which depends on iteration number t).



Shows the result for a seed with various scaling factor 'h' parameters. The shadows are generated using the code and the solid lines utilise the Savitzky-Golay Filter[2] (MSE line) for smoothness.

### Diamond Square

This method begins with a 2D grid and creates a matrix of heightmap values. These essentially are the z coordinates used for generating a 3D surface. Starting with a 2D array, initialise the 4 corners with some random value. Firstly, a diamond step takes an average of the 4 nearest corners and adds some random noise - think of this as diagonal. Secondly, a square step takes an average of the nearest horizontal and vertical points (an average of 3 if an edge) with additional noise. These are repeated until the array is full, decreasing the range of randomness with each iteration.



[3]A visualisation of the Diamond-Square process.


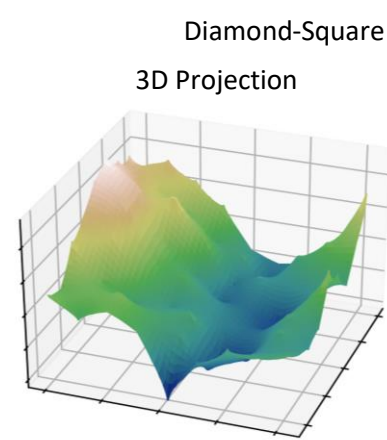
After finding patterns in each iteration (one diamond step and one square step) we have identified formulas to make an efficient algorithm function successfully. Attached are figures of the pseudo-code and an example of how we use loops in python to target each cell, depending on the current iteration t.

```
"The following script can be directly imported for easy use within scripts"
"we have numbered the order in which the algorithm computes on the left"

In short, the function is called with the following parameters:
- n: calculates the dimension & iterations (default n=8, recommended 10-12)
- seed: the iteration number (seed=0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84, 85, 86, 87, 88, 89, 90, 91, 92, 93, 94, 95, 96, 97, 98, 99, 100)
- r: is the initial range of random integer values (default n=8)

def diamondSquare(n, seed=0, r=8):
    # Create a 2D array of size n x n
    # Initialize the 4 corners with random values
    # Iterate until the array is full
    # Use diamond and square steps to update values
    # Return the final heightmap values
```


Diamond Square pseudo-code



$$t = 1, 2, \dots, n \quad m_t = \frac{d-1}{t} \quad e_t = \frac{m_t}{2} \quad a_{ds} = [0, 1, \dots, \sqrt{C_{ds}}] \quad b_{ds} = [0, 1, \dots, \sqrt{C_{ds}}]$$
$$d = 2^n + 1 \quad C_{ds} = 2^{2(t-1)}$$
$$[a_{ds}, m_t][b_{ds}, m_t] \quad [a_{ds}, m_t][b_{ds} + 1, m_t]$$
$$[a_{ds}, m_t + e_t][b_{ds}, m_t + e_t] \quad [(a_{ds} + 1), m_t][b_{ds}, m_t]$$

### Perlin Noise

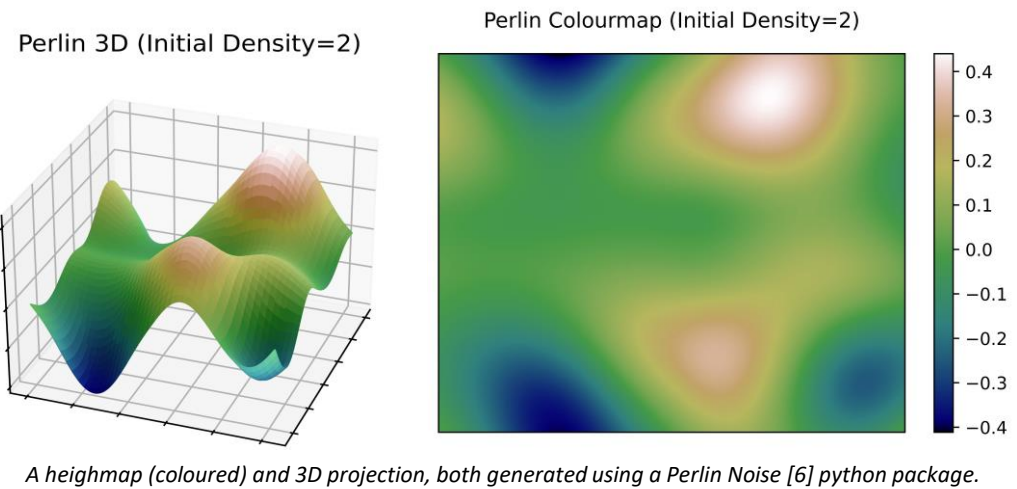
This method was developed by Ken Perlin [4] and creates a heightmap for any dimension. Here, we look at 2D to 3D. The process is simply demonstrated as follows, from A > B > C:



[5]Perlin Noise Theory.

1. Create a 2D grid of cells and assign each intersection with a unit vector of gradient between -1 and 1.
2. For every pixel in the whole grid, calculate the displacement vector (d) of said point with its nearest cell corner.
3. Calculate the dot product of all these d vectors and the corresponding corner vectors. This is shown in image B (colour coordinated).
4. Repeat this process but to the other 3 corners of each cell to get 4 total sets of gradients.
5. Combine the 4 sets of the gradients (dot product values) together using a "smoothstep function" to get to the heightmap C.
6. Numerous versions of these are then combined to increase clarity. Known as octaves, the density of the initial grid increases, these are added together with a decreased weight / transparency.

The final result produces incredible, smooth terrain.



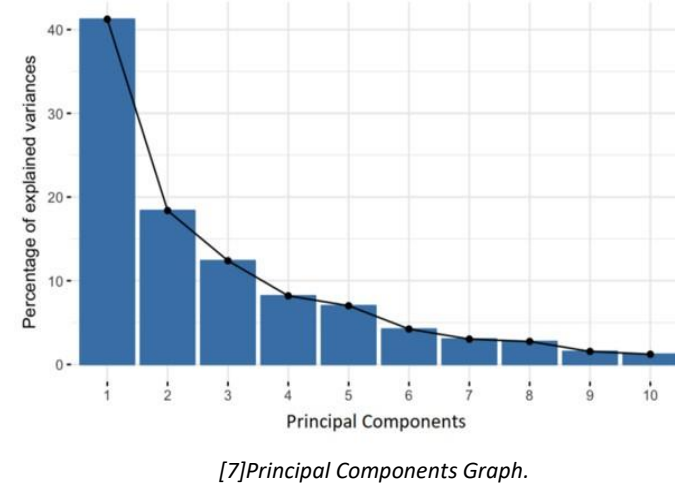
A heighmap (coloured) and 3D projection, both generated using a Perlin Noise [6] python package.

## Principal Component Analysis

Explanation:

Principal Component Analysis (PCA) is a statistical method used for dimensionality reduction in data by transforming variables into a new set of uncorrelated components, (called principal values). The most significant information is captured, while minimizing loss of data variance.

Here you can see the impact the number of principal values has on the amount of variance explained. As the number of principal values are increased, they start to explain sequentially less of the variance. Therefore, we must determine how many of these values we want.



[7]Principal Components Graph.

We plan to apply PCA to analyse satellite-captured LiDAR scans, converting elevation data into a Python array, and using the PCA algorithm to select the optimal number of principal components for the analysis, as demonstrated locally in Keswick.

Procedure:

1. Standardise the range.

$$z = (x - \mu) / \sigma$$

2. Compute the covariance matrix, then remove data points with high similarity to preserve information and remove data.

$$Cov(X, Y) = \sum_{i=1}^n \frac{(X_i - \bar{X})(Y_i - \bar{Y})}{n - 1}$$

$$\begin{bmatrix} Cov(x, x) & Cov(x, y) & Cov(x, z) \\ Cov(y, x) & Cov(y, y) & Cov(y, z) \\ Cov(z, x) & Cov(z, y) & Cov(z, z) \end{bmatrix}$$




[8]LiDAR data - Keswick

## Autoencoders

Process:

- An autoencoder is designed to reduce the dimensionality of data while minimising the loss of the key information.
- It does this by passing the data through multiple encoding layers which reduces the dimensionality of the data.
- The data is compressed as it passes through the bottleneck allowing only the most important information to pass through.
- The information is then reconstructed as close as possible back to its original state by using multiple decoding layers.
- At the bottleneck layer, the trained network will have compressed our terrain data into a low-dimensional latent space.
- By then randomly sampling from this space, we can get the decoder to produce completely new terrain.
- This new terrain will be similar to terrain the network was trained on and in addition very similar to particular terrains mapped to that region of the latent space.



[9]Autoencoder Theory - a visual representation of the neural network structure.

$$MSE = \frac{1}{n} \sum_{i=1}^n \iota(x_i, x'_i)$$

[10]Mean Square Error loss function used in backpropagation.

## Conclusion and Next Steps

Conclusion:

In summary, we have looked at the different methods used for terrain generation and carried out these methods to produce our own examples of how each method can generate terrain. We have explained how Principal Component Analysis can reduce the dimensionality of data while capturing the most important information and how this process is used in an important part of an autoencoder.

Next Steps:

- We are planning on writing code to convert the LiDAR scans into an array within Python, and then use a pre-existing library for Principal Component Analysis within Sci-Kit to reduce this data.
- We aim to train an auto-encoder on both the generated and real data by importing a machine learning library PyTorch, then creating new terrains from it.

## References

- [1] Minecraft Landscape - [https://www.reddit.com/r/Minecraft/comments/r25c15/mountain\\_landscape\\_i\\_made\\_over\\_the\\_last\\_few\\_days](https://www.reddit.com/r/Minecraft/comments/r25c15/mountain_landscape_i_made_over_the_last_few_days)
- [2] Savitzky-Golay Filter - <https://www.mathworks.com/help/signal/ref/sgolayfilt.html>
- [3] Diamond Square Theory - [https://en.wikipedia.org/wiki/Diamond-square\\_algorithm](https://en.wikipedia.org/wiki/Diamond-square_algorithm)
- [4] Ken Perlin's Paper on Perlin Noise: An Image Synthesiser - ACM SIGGRAPH Computer Graphics, 19(3):287-296, 1985
- [5] Perlin Noise Theory - [https://en.wikipedia.org/wiki/Perlin\\_noise](https://en.wikipedia.org/wiki/Perlin_noise)
- [6] Perlin Noise Python Package - <https://pypi.org/project/perlin-noise>
- [7] Principal Components Graph - [https://link.springer.com/chapter/10.1007/978-3-031-04420-5\\_4](https://link.springer.com/chapter/10.1007/978-3-031-04420-5_4)
- [8] LiDAR Data - <https://www.data.gov.uk/dataset/f0db0249-f17b-4036-9e65-309148c97ce4/national-lidar-programme>
- [9] Autoencoder Theory - <https://towardsdatascience.com/applied-deep-learning-part-3-autoencoders-1c083af4d798>
- [10] Mean Square Error loss function - [https://www.researchgate.net/publication/339480272\\_An\\_Overview\\_of\\_Deep\\_Learning\\_Architecture\\_of\\_Deep\\_Neural\\_Networks\\_and\\_Autoencoders](https://www.researchgate.net/publication/339480272_An_Overview_of_Deep_Learning_Architecture_of_Deep_Neural_Networks_and_Autoencoders)